Survival Ensembles

Torsten Hothorn[1,*], Peter Bühlmann[2], Sandrine Dudoit[3],
Annette Molinaro[4] and Mark J. van der Laan[3]

[1]Institut für Medizininformatik, Biometrie und Epidemiologie
Friedrich-Alexander-Universität Erlangen-Nürnberg
Waldstraße 6, D-91054 Erlangen, Germany
Tel: ++49–9131–8522707
Fax: ++49–9131–8525740
Torsten.Hothorn@R-project.org

[2]Seminar für Statistik, ETH Zürich, CH-8032 Zürich, Switzerland
buhlmann@stat.math.ethz.ch

[3]Division of Biostatistics, University of California, Berkeley
140 Earl Warren Hall, #7360, Berkeley, CA 94720-7360, USA
sandrine@stat.Berkeley.EDU
laan@stat.Berkeley.EDU

[4]Division of Biostatistics, Epidemiology and Public Health
Yale University School of Medicine, 206 LEPH
60 College Street PO Box 208034, New Haven CT 06520-8034
annette.molinaro@yale.edu

1 Illustrations and Applications

This document reproduces the data analyses presented in Hothorn et al. (2006).
For a description of the theory behind applications shown here we refer to the
original manuscript.

1.1 Acute myeloid leukemia

Data preprocessing Compute IPC weights, define risk score and set up
learning sample:

```
R> AMLw <- IPCweights(Surv(clinical$time, clinical$event))
R> risk <- rep(0, nrow(clinical))
R> rlev <- levels(clinical[, "Cytogenetic.group"])
```

```
R> risk[clinical[, "Cytogenetic.group"] %in% rlev[c(7,
        8, 4)]] <- "low"
R> risk[clinical[, "Cytogenetic.group"] %in% rlev[c(5,
        9)]] <- "intermediate"
R> risk[clinical[, "Cytogenetic.group"] %in% rlev[-c(4,
        5, 7, 8, 9)]] <- "high"
R> risk <- as.factor(risk)
R> AMLlearn <- cbind(clinical[, c("time", "Sex",
        "Age", "LDH", "WBC", "FLT3.aberration.", "MLL.PTD",
        "Tx.Group.")], risk = risk, iexpressions[,
        colnames(iexpressions) %in% selgenes[["Clone.ID"]]])
R> cc <- complete.cases(AMLlearn)
R> AMLlearn <- AMLlearn[AMLw > 0 & cc, ]
R> AMLw <- AMLw[AMLw > 0 & cc]
```

Model fitting Fit random forest for censored data

```
R> ctrl <- cforest_control(mincriterion = 0.1, mtry = 5,
        minsplit = 5, ntree = 250)
R> AMLrf <- cforest(I(log(time)) ~ ., data = AMLlearn,
        control = ctrl, weights = AMLw)
```

and L_2Boosting for censored data

```
R> AMLl2b <- glmboost(I(log(time)) ~ ., data = AMLlearn,
        weights = AMLw, control = boost_control(mstop = 5000))
```

   Compute fitted values

```
R> AMLl2b <- AMLl2b[mstop(aic)]
R> cAML <- coef(AMLl2b)
R> cAML[abs(cAML) > 0]

    (Intercept)              Age               WBC
     0.03094981       0.00854937       -0.00364371
       MLL.PTDyes    Tx.Group.AUTO      Tx.Group.IC
     -0.50709786       0.90185340       0.04037578
    Tx.Group.Ind riskintermediate    `IMAGE:145643`
     -1.86134842       0.11825619       0.19788355
 `IMAGE:2542486`   `IMAGE:345601`   `IMAGE:377560`
     0.00442375       0.02935101       0.11000322
  `IMAGE:428782`  `IMAGE:2043415`  `IMAGE:1584563`
     0.01010658       0.05911671       -0.17883619
  `IMAGE:347035`   `IMAGE:262695`   `IMAGE:950479`
     -0.03307600       0.00080156       0.09049309
  `IMAGE:898305`  `IMAGE:1472689`   `IMAGE:150702`
     0.00523016       0.03498572       0.01367553
 `IMAGE:1526826`    `IMAGE:66507`   `IMAGE:786302`
```

```
R> plot(aic <- AIC(AMLl2b))
```


Figure 1: AIC criterion for AML data.

```
      -0.01805326         0.00399127         0.08941300
   `IMAGE:243614`      `IMAGE:417884`    `IMAGE:1592006`
      -0.05776062        -0.04890054        -0.02269622
  `IMAGE:1917063`      `IMAGE:884333`     `IMAGE:133273`
      -0.06536720         0.04189990         0.06594787
   `IMAGE:950888`      `IMAGE:809533`      `IMAGE:49389`
       0.02027810        -0.15986981         0.06352703
   `IMAGE:789357`      `IMAGE:142139`    `IMAGE:1558053`
      -0.01252187         0.00089307         0.07795515
   `IMAGE:856174`      `IMAGE:504421`     `IMAGE:435036`
       0.01115234         0.06861766         0.06094620
   `IMAGE:491751`      `IMAGE:782835`      `IMAGE:52930`
       0.04336285        -0.17924185        -0.03503330
  `IMAGE:2545705`      `IMAGE:756405`     `IMAGE:502664`
      -0.09886616         0.07713650         0.03620466
   `IMAGE:129032`     `IMAGE:1610168`     `IMAGE:327676`
      -0.31322459         0.01260374        -0.02117310
    `IMAGE:69002`      `IMAGE:121551`    `IMAGE:2019101`
      -0.41671336        -0.08107446        -0.06531175
  `IMAGE:1456160`      `IMAGE:430318`    `IMAGE:2566064`
      -0.10208684        -0.07297586         0.06126683
    `IMAGE:74537`     `IMAGE:1606557`     `IMAGE:306812`
       0.04523784         0.14243526         0.03504441
   `IMAGE:565083`      `IMAGE:843028`      `IMAGE:68794`
       0.29555347         0.05619983         0.23722775
   `IMAGE:488505`      `IMAGE:167205`     `IMAGE:291756`
       0.33464829         0.00217136         0.04973319
   `IMAGE:810801`     `IMAGE:1702742`     `IMAGE:380462`
       0.08725523        -0.04428190        -0.13182519
   `IMAGE:154472`      `IMAGE:302540`     `IMAGE:135221`
      -0.24723347         0.17175129        -0.01972168
  `IMAGE:1567220`      `IMAGE:594630`
       0.02473376        -0.07396882

R> AMLprf <- predict(AMLrf, newdata = AMLlearn)
R> AMLpb <- predict(AMLl2b, newdata = AMLlearn)
```

1.2 Node-positive breast cancer

Data preprocessing Compute IPC weights and set up learning sample:

```
R> data("GBSG2", package = "ipred")
R> GBSG2w <- IPCweights(Surv(GBSG2$time, GBSG2$cens))
R> GBSG2learn <- cbind(GBSG2[, -which(names(GBSG2) %in%
        c("time", "cens"))], ltime = log(GBSG2$time))
R> n <- nrow(GBSG2learn)
```
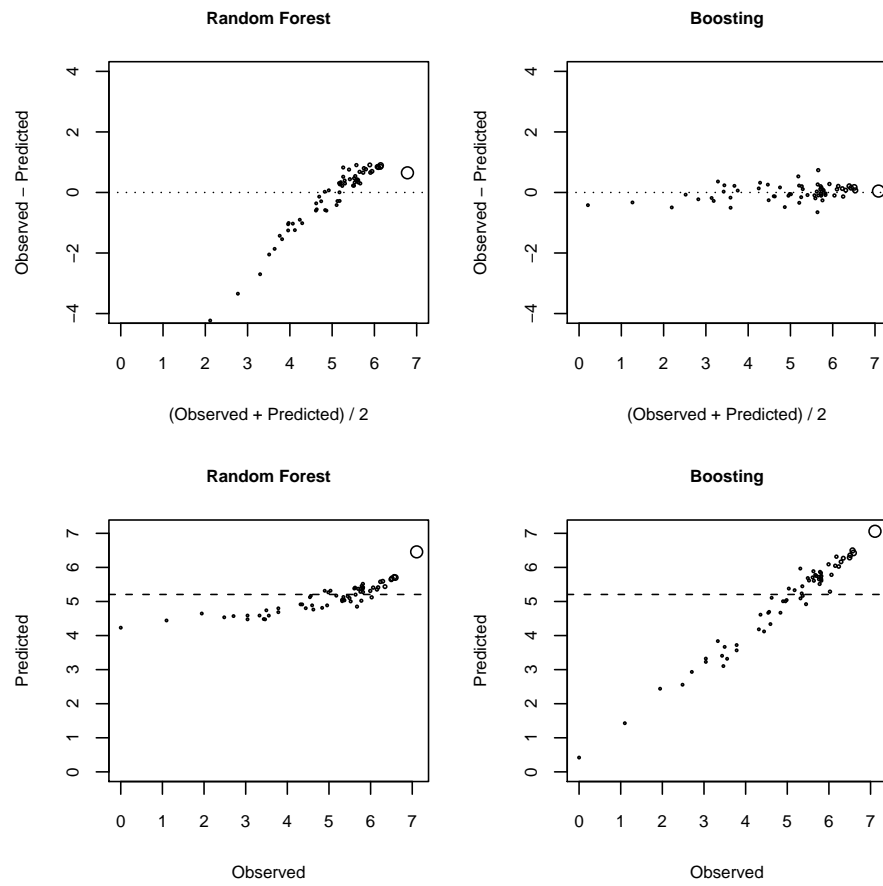
Figure 2: AML data: Reproduction of Figure 1.

Model fitting

```
R> LMmod <- lm(ltime ~ ., data = GBSG2learn, weights = GBSG2w)
R> LMerisk <- sum((GBSG2learn$ltime - predict(LMmod))^2 *
        GBSG2w)/n
R> TRmod <- rpart(ltime ~ ., data = GBSG2learn, weights = GBSG2w)
R> TRerisk <- sum((GBSG2learn$ltime - predict(TRmod))^2 *
        GBSG2w)/n
R> ctrl <- cforest_control(mincriterion = qnorm(0.95),
        mtry = 5, minsplit = 5, ntree = 100)
R> RFmod <- cforest(ltime ~ ., data = GBSG2learn,
        weights = GBSG2w, control = ctrl)
R> L2Bmod <- glmboost(ltime ~ ., data = GBSG2learn,
        weights = GBSG2w, control = boost_control(mstop = 250))
R> L2BHubermod <- glmboost(ltime ~ ., data = GBSG2learn,
        weights = GBSG2w, family = Huber(d = log(2)))
```

Compute fitted values:

```
R> GBSG2Hp <- predict(L2BHubermod, newdata = GBSG2learn)
R> L2Berisk <- sum((GBSG2learn$ltime - predict(L2Bmod,
        newdata = GBSG2learn))^2 * GBSG2w)/n
R> RFerisk <- sum((GBSG2learn$ltime - predict(RFmod,
        newdata = GBSG2learn))^2 * GBSG2w)/n
```

```
R> plot(aic <- AIC(L2Bmod))
```


Figure 3: AIC criterion for GBSG2 data.

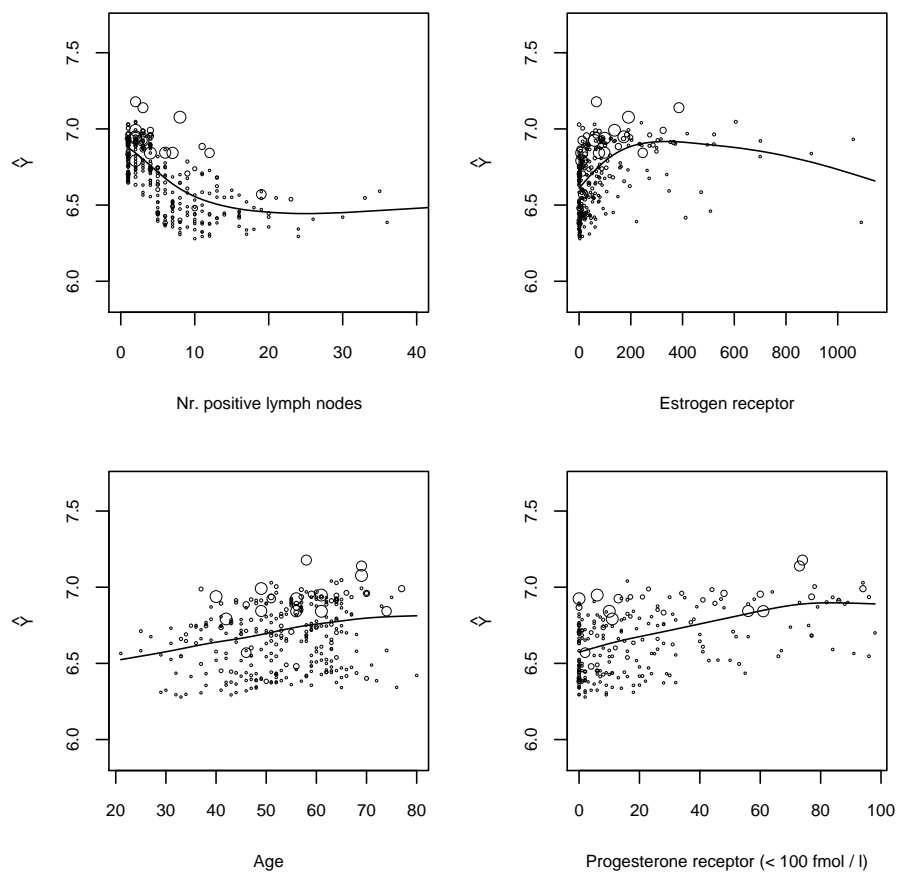Figure 4: GBSG-2 data: Reproduction of Figure 3.

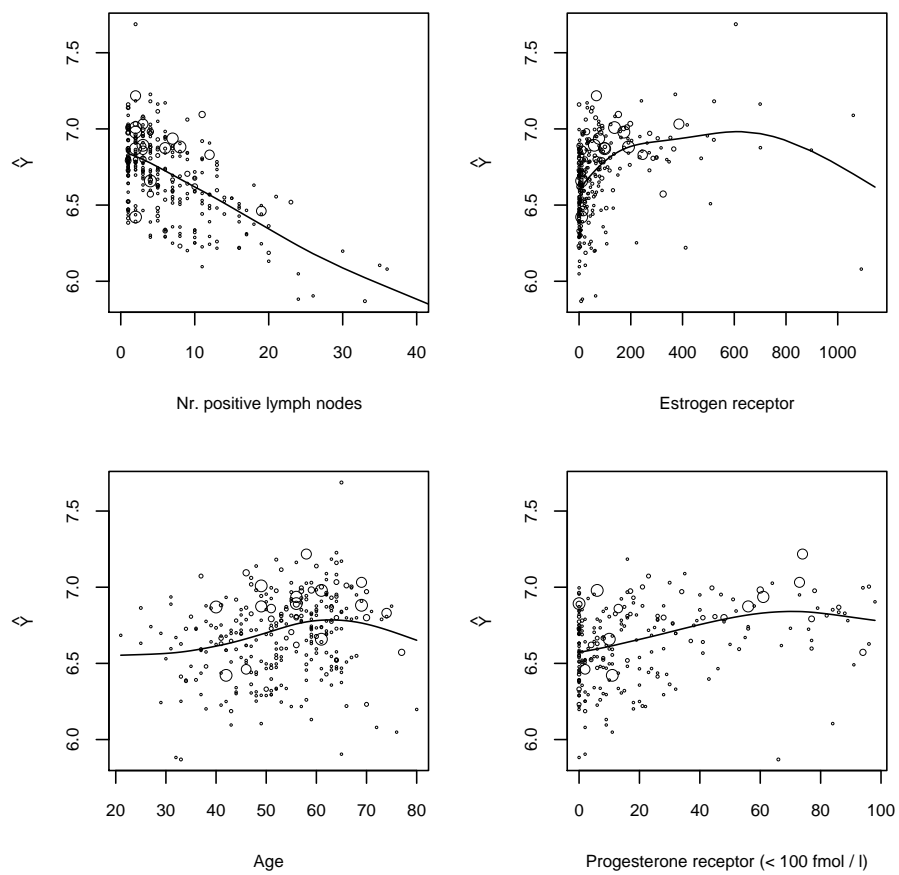Figure 5: GBSG-2 data: Reproduction of Figure 5.
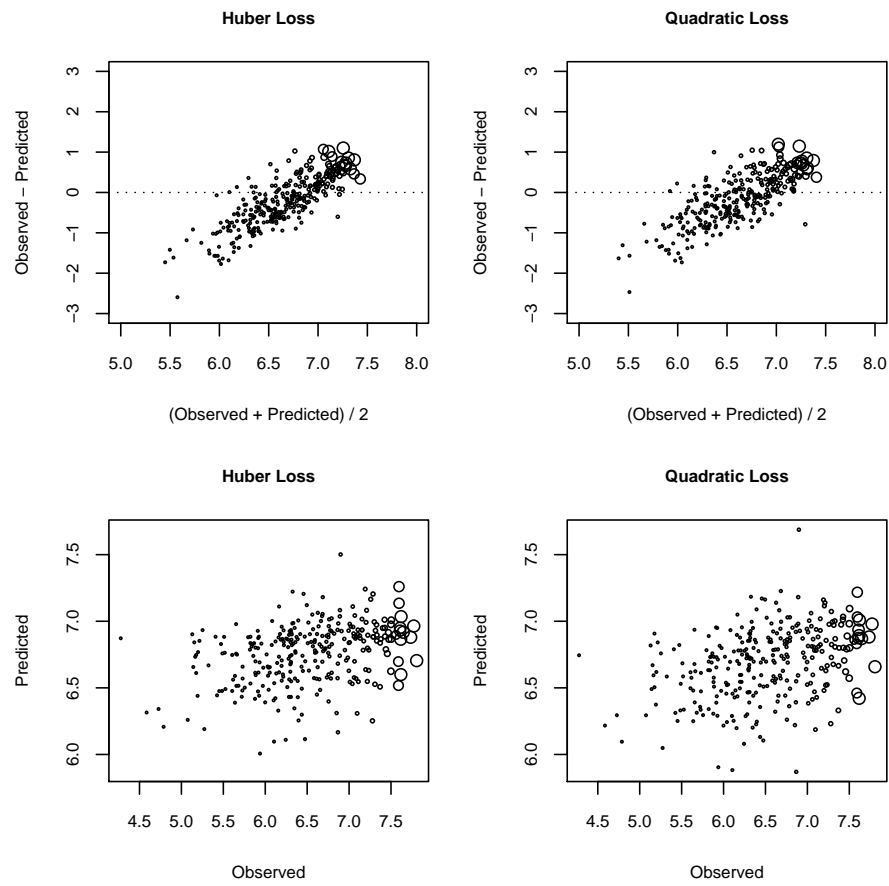
Figure 6: GBSG-2 data: Reproduction of Figure 6.

Figure 7: GBSG-2 data: Reproduction of Figure 7.

References

T. Hothorn, P. Bühlmann, S. Dudoit, A. Molinaro, and M. van der Laan. Survival ensembles. *Biostatistics*, 2006. in press.