Survival Ensembles

Torsten Hothorn[1,*], Peter Bühlmann[2], Sandrine Dudoit[3],
Annette Molinaro[4] and Mark J. van der Laan[3]

[1]Institut für Statistik
Ludwig-Maximilians-Universität München
Ludwigstraße 33, D-80539 München, Germany
Tel: ++49–9131–8522707
Fax: ++49–9131–8525740
Torsten.Hothorn@R-project.org

[2]Seminar für Statistik, ETH Zürich, CH-8032 Zürich, Switzerland
buhlmann@stat.math.ethz.ch

[3]Division of Biostatistics, University of California, Berkeley
140 Earl Warren Hall, #7360, Berkeley, CA 94720-7360, USA
sandrine@stat.Berkeley.EDU
laan@stat.Berkeley.EDU

[4]Division of Biostatistics, Epidemiology and Public Health
Yale University School of Medicine, 206 LEPH
60 College Street PO Box 208034, New Haven CT 06520-8034
annette.molinaro@yale.edu

1 Illustrations and Applications

This document reproduces the data analyses presented in Hothorn et al. (2006). For a description of the theory behind applications shown here we refer to the original manuscript. The results differ slightly due to technical changes or bug-fixes in **mboost** that have been implemented after the paper was printed.

1.1 Acute myeloid leukemia

Data preprocessing Compute IPC weights, define risk score and set up learning sample:

```
R> AMLw <- IPCweights(Surv(clinical$time, clinical$event))
R> risk <- rep(0, nrow(clinical))
```

```
R> rlev <- levels(clinical[, "Cytogenetic.group"])
R> risk[clinical[, "Cytogenetic.group"] %in% rlev[c(7,
        8, 4)]] <- "low"
R> risk[clinical[, "Cytogenetic.group"] %in% rlev[c(5,
        9)]] <- "intermediate"
R> risk[clinical[, "Cytogenetic.group"] %in% rlev[-c(4,
        5, 7, 8, 9)]] <- "high"
R> risk <- as.factor(risk)
R> AMLlearn <- cbind(clinical[, c("time", "Sex",
        "Age", "LDH", "WBC", "FLT3.aberration.", "MLL.PTD",
        "Tx.Group.")], risk = risk, iexpressions[,
        colnames(iexpressions) %in% selgenes[["Clone.ID"]]])
R> cc <- complete.cases(AMLlearn)
R> AMLlearn <- AMLlearn[AMLw > 0 & cc, ]
R> AMLw <- AMLw[AMLw > 0 & cc]
```

Model fitting   Fit random forest for censored data

```
R> ctrl <- cforest_control(mincriterion = 0.1, mtry = 5,
        minsplit = 5, ntree = 250)
R> AMLrf <- cforest(I(log(time)) ~ ., data = AMLlearn,
        control = ctrl, weights = AMLw)
```

and L_2Boosting for censored data

```
R> AMLl2b <- glmboost(I(log(time)) ~ ., data = AMLlearn,
        weights = AMLw, control = boost_control(mstop = 5000))
```

Compute fitted values

```
R> AMLl2b <- AMLl2b[mstop(aic)]
R> cAML <- coef(AMLl2b)
R> cAML[abs(cAML) > 0]

    (Intercept)              Age              WBC
     0.03094981       0.00854937      -0.00364371
      MLL.PTDyes    Tx.Group.AUTO     Tx.Group.IC
     -0.50709786       0.90185340       0.04037578
    Tx.Group.Ind riskintermediate  `IMAGE:145643`
     -1.86134842       0.11825619       0.19788355
 `IMAGE:2542486`   `IMAGE:345601`   `IMAGE:377560`
      0.00442375       0.02935101       0.11000322
  `IMAGE:428782`  `IMAGE:2043415`  `IMAGE:1584563`
      0.01010658       0.05911671      -0.17883619
  `IMAGE:347035`   `IMAGE:262695`   `IMAGE:950479`
     -0.03307600       0.00080156       0.09049309
  `IMAGE:898305`  `IMAGE:1472689`   `IMAGE:150702`
      0.00523016       0.03498572       0.01367553
 `IMAGE:1526826`    `IMAGE:66507`   `IMAGE:786302`
     -0.01805326       0.00399127       0.08941300
  `IMAGE:243614`   `IMAGE:417884`  `IMAGE:1592006`
```

```
R> plot(aic <- AIC(AMLl2b))
```


Figure 1: AIC criterion for AML data.

```
      -0.05776062       -0.04890054       -0.02269622
 `IMAGE:1917063`    `IMAGE:884333`    `IMAGE:133273`
      -0.06536720        0.04189990        0.06594787
  `IMAGE:950888`    `IMAGE:809533`     `IMAGE:49389`
       0.02027810       -0.15986981        0.06352703
  `IMAGE:789357`    `IMAGE:142139`   `IMAGE:1558053`
      -0.01252187        0.00089307        0.07795515
  `IMAGE:856174`    `IMAGE:504421`    `IMAGE:435036`
       0.01115234        0.06861766        0.06094620
  `IMAGE:491751`    `IMAGE:782835`     `IMAGE:52930`
       0.04336285       -0.17924185       -0.03503330
 `IMAGE:2545705`    `IMAGE:756405`    `IMAGE:502664`
      -0.09886616        0.07713650        0.03620466
  `IMAGE:129032`   `IMAGE:1610168`    `IMAGE:327676`
      -0.31322459        0.01260374       -0.02117310
   `IMAGE:69002`    `IMAGE:121551`   `IMAGE:2019101`
      -0.41671336       -0.08107446       -0.06531175
 `IMAGE:1456160`    `IMAGE:430318`   `IMAGE:2566064`
      -0.10208684       -0.07297586        0.06126683
   `IMAGE:74537`   `IMAGE:1606557`    `IMAGE:306812`
       0.04523784        0.14243526        0.03504441
  `IMAGE:565083`    `IMAGE:843028`     `IMAGE:68794`
       0.29555347        0.05619983        0.23722775
  `IMAGE:488505`    `IMAGE:167205`    `IMAGE:291756`
       0.33464829        0.00217136        0.04973319
  `IMAGE:810801`   `IMAGE:1702742`    `IMAGE:380462`
       0.08725523       -0.04428190       -0.13182519
  `IMAGE:154472`    `IMAGE:302540`    `IMAGE:135221`
      -0.24723347        0.17175129       -0.01972168
 `IMAGE:1567220`    `IMAGE:594630`
       0.02473376       -0.07396882
```

```
R> AMLprf <- predict(AMLrf, newdata = AMLlearn)
R> AMLpb <- predict(AMLl2b, newdata = AMLlearn)
```

1.2   Node-positive breast cancer

Data preprocessing   Compute IPC weights and set up learning sample:

```
R> data("GBSG2", package = "ipred")
R> GBSG2w <- IPCweights(Surv(GBSG2$time, GBSG2$cens))
R> GBSG2learn <- cbind(GBSG2[, -which(names(GBSG2) %in%
        c("time", "cens"))], ltime = log(GBSG2$time))
R> n <- nrow(GBSG2learn)
```

Model fitting

```
R> LMmod <- lm(ltime ~ ., data = GBSG2learn, weights = GBSG2w)
R> LMerisk <- sum((GBSG2learn$ltime - predict(LMmod))^2 *
        GBSG2w)/n
R> TRmod <- rpart(ltime ~ ., data = GBSG2learn, weights = GBSG2w)
R> TRerisk <- sum((GBSG2learn$ltime - predict(TRmod))^2 *
```
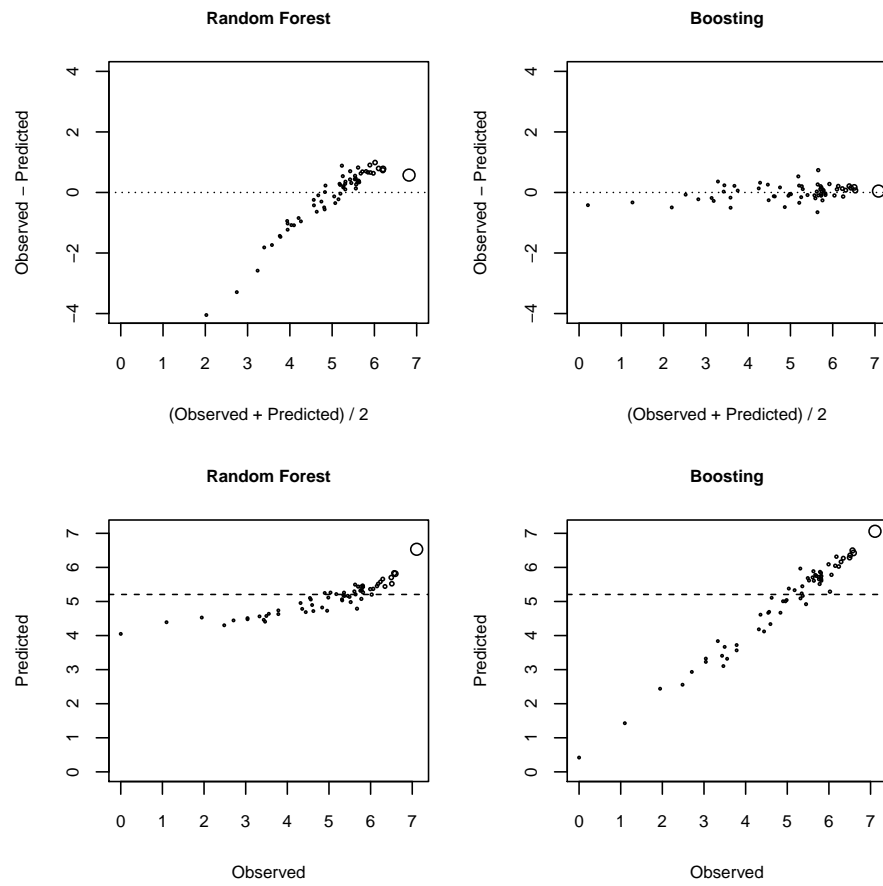
Figure 2: AML data: Reproduction of Figure 1.

```
             GBSG2w)/n
R> ctrl <- cforest_control(mincriterion = qnorm(0.95),
          mtry = 5, minsplit = 5, ntree = 100)
R> RFmod <- cforest(ltime ~ ., data = GBSG2learn,
          weights = GBSG2w, control = ctrl)
R> L2Bmod <- glmboost(ltime ~ ., data = GBSG2learn,
          weights = GBSG2w, control = boost_control(mstop = 250))
R> L2BHubermod <- glmboost(ltime ~ ., data = GBSG2learn,
          weights = GBSG2w, family = Huber(d = log(2)))
```

Compute fitted values:

```
R> GBSG2Hp <- predict(L2BHubermod, newdata = GBSG2learn)
R> L2Berisk <- sum((GBSG2learn$ltime - predict(L2Bmod,
          newdata = GBSG2learn))^2 * GBSG2w)/n
R> RFerisk <- sum((GBSG2learn$ltime - predict(RFmod,
          newdata = GBSG2learn))^2 * GBSG2w)/n
```

```
R> plot(aic <- AIC(L2Bmod))
```


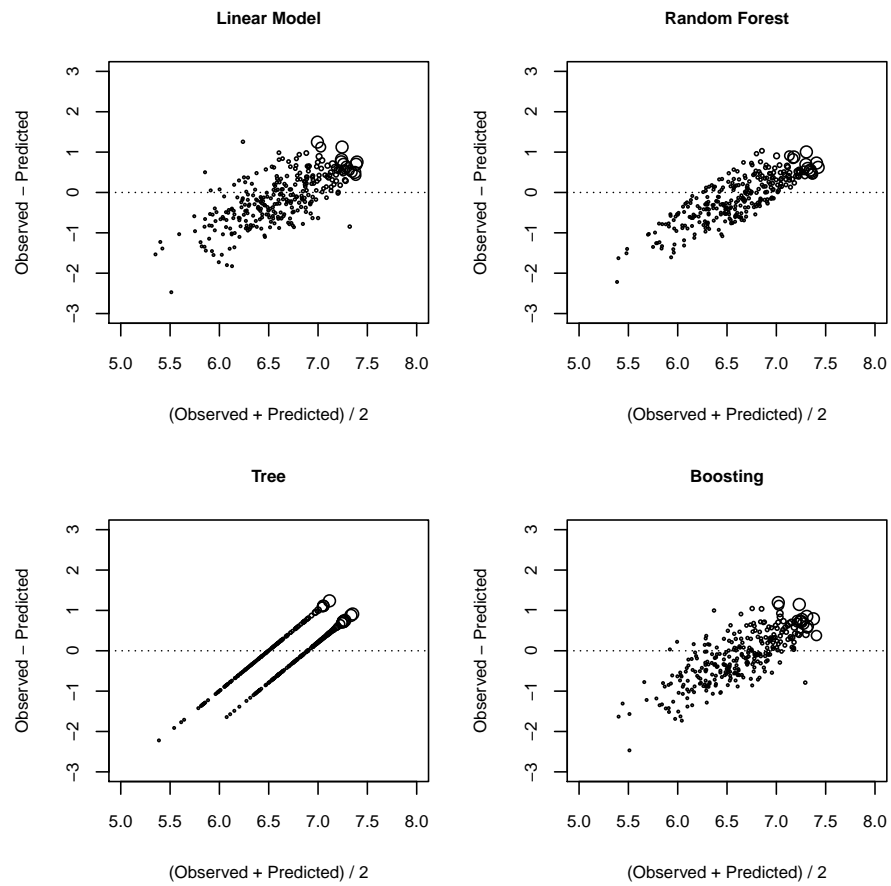Figure 3: AIC criterion for GBSG2 data.

Figure 4: GBSG-2 data: Reproduction of Figure 3.

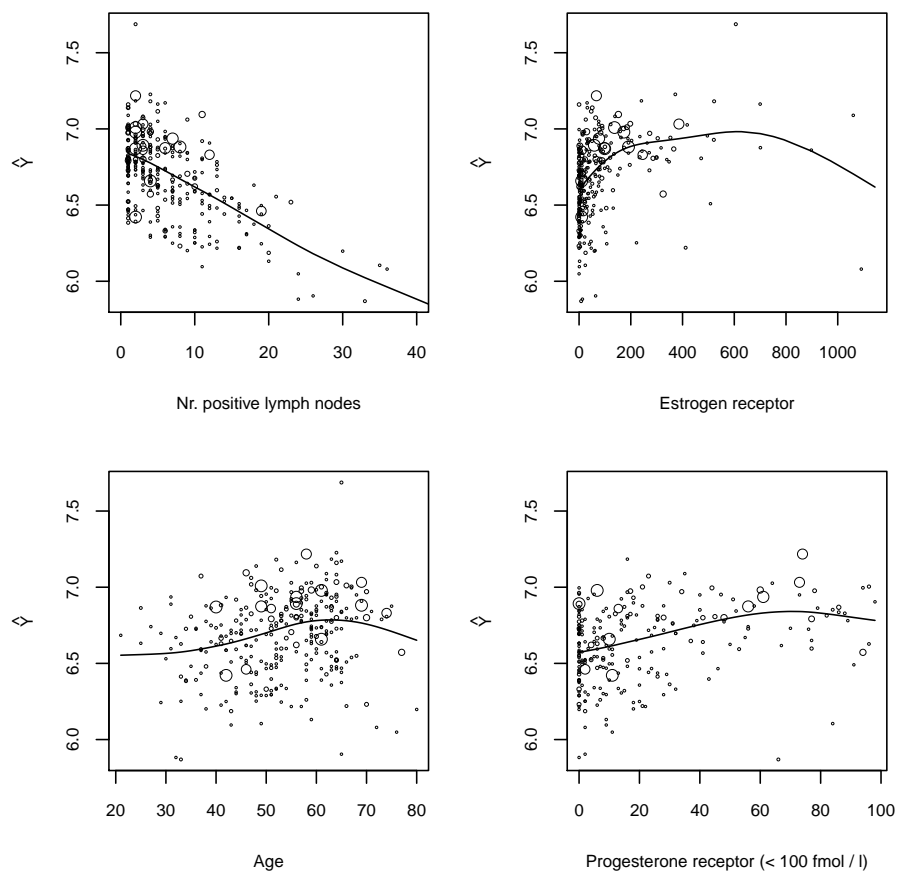Figure 5: GBSG-2 data: Reproduction of Figure 5.
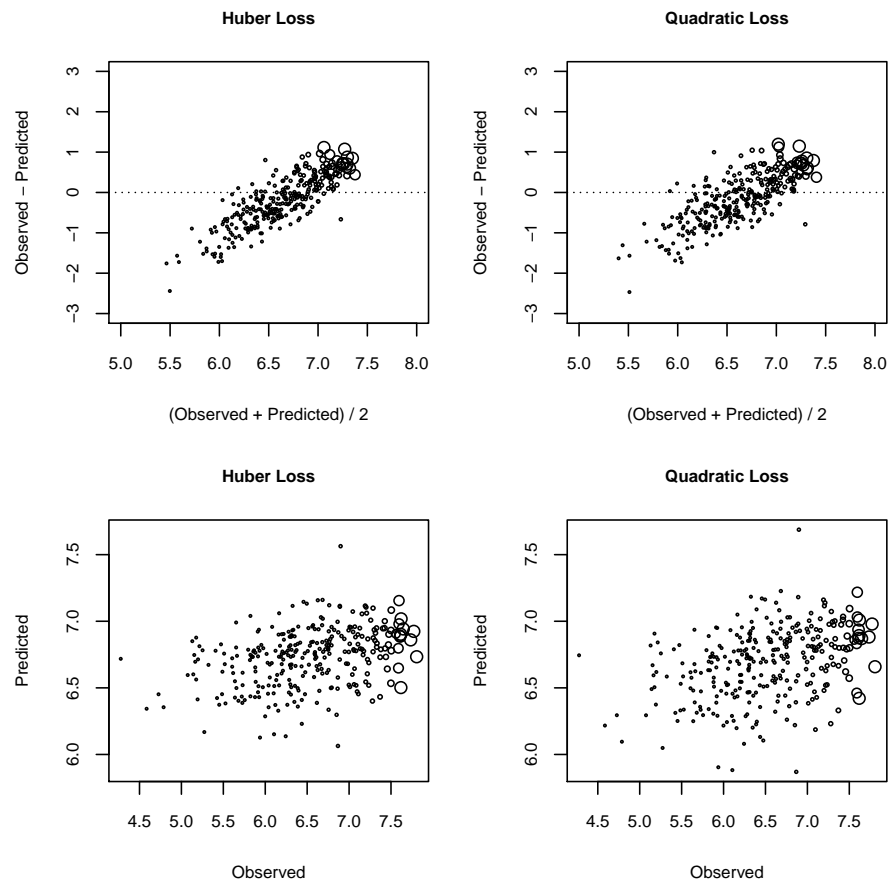
Figure 6: GBSG-2 data: Reproduction of Figure 6.

Figure 7: GBSG-2 data: Reproduction of Figure 7.

References

T. Hothorn, P. Bühlmann, S. Dudoit, A. Molinaro, and M. van der Laan. Survival ensembles. *Biostatistics*, 7:355–373, 2006.